

[CSED233-01] Data Structure
2-3 Tree, B-Tree

Jaesik Park

POSTECH

B

Balanced Search Trees

- **Memory-based** search trees
 - Balanced BST
 - **AVL** (Adelson-Velskii & Landis) trees
 - Red-black trees
 - Splay trees, ...
 - Balanced multi-way search tree
 - **2-3, 2-3-4 trees (B-trees)**
- **Disk-based** search trees
 - Balanced multi-way search trees
 - **B-trees (B+, B*)**
 - Prefix B-trees


2-3 Tree: Shape Property

- A tree with the following *shape properties*:
 - Each internal node has 2 or 3 children
 - All leaves are at the same level (→ Perfectly height-balanced)
- Internal nodes
 - Store one or two key values
 - Placeholders to guide the search
- Leaf nodes
 - All *actual records* - (*key*, *info*) pairs are stored at the leaves

2-3 Tree: Search Tree Property


- The 2-3 tree has a *search tree property*:

- All keys in left subtree are smaller than first key
- All keys in middle subtree are greater than or equal to first key, and smaller than second key
- All keys in right subtree are greater than or equal to second key



2-3 Tree: Example


- Actual records are in the *leaves*
 - B⁺-tree of order 3



- The keys are in the leaves
 - linearly ordered from left to right

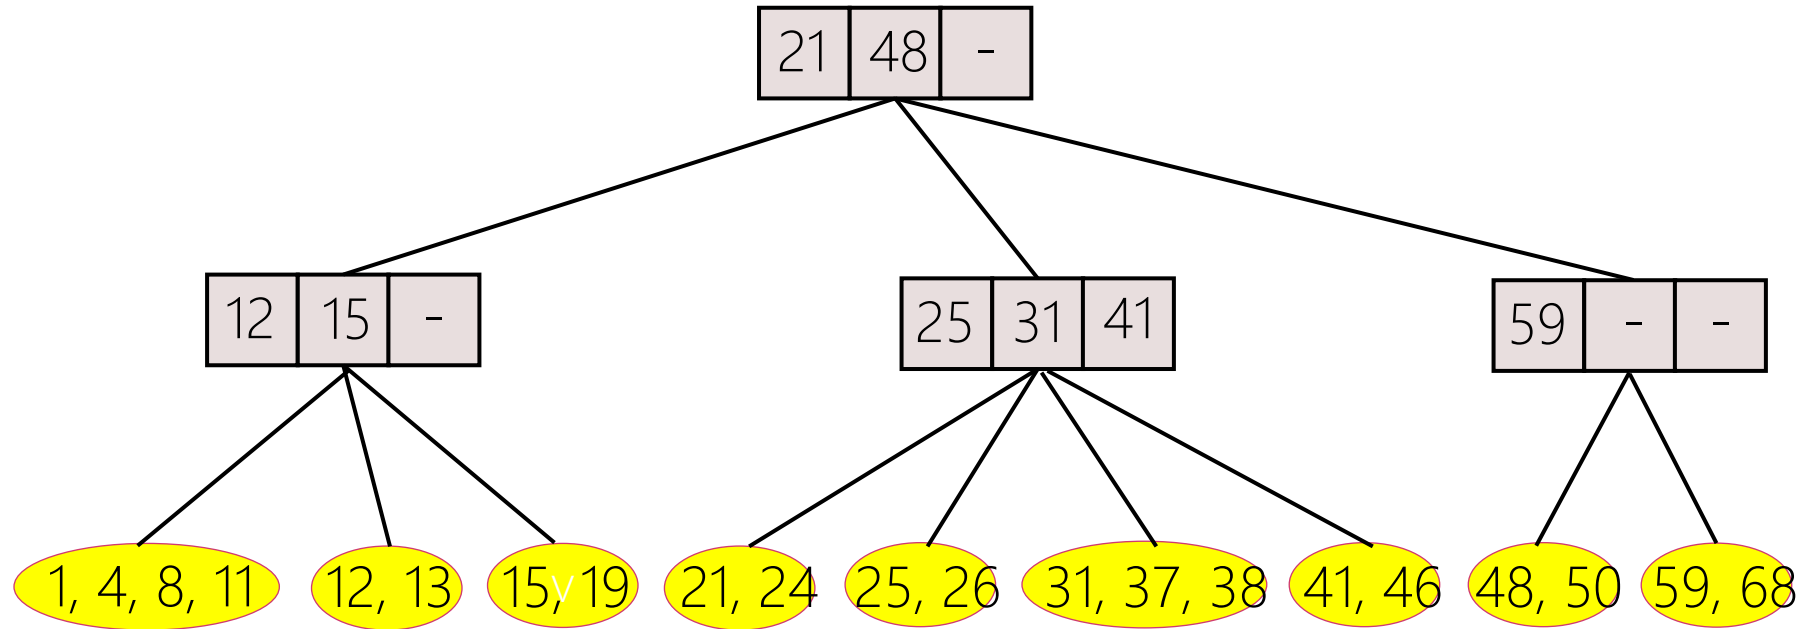
2-3 Tree: Alternative Structure

- Actual records are stored in both *leaves* and *internal nodes*
 - B-tree of order 3



2-3-4 Tree: Example

- B⁺-tree of order 4




- A leaf node in a B⁺-tree of order m
 - may store more or less than m (generally, between $\lceil m/2 \rceil$ and m) actual records

2-3 Tree: Height Bound

- 2-3 tree of h levels (of $h-1$ height) has
 - At least 2^{h-1} leaves (since each internal node has at least 2 children)
 - At most 3^{h-1} leaves (since each internal node has at most 3 children)
 - Let n : the # of leaves in a 2-3 tree
 - $2^{h-1} \leq n \leq 3^{h-1}$
 - $(\log_3 n + 1) \leq h \leq (\log_2 n + 1)$
- ➔ Path length (height) = ? $O(\log n)$

2-3 Tree: Search

- $search(x, tree)$
 - If $x < k_1$, $search(x, left_subtree)$
 - If $k_1 < x < k_2$, $search(x, middle_subtree)$
 - If $x > k_2$, $search(x, right_subtree)$
 - If $x = k_1$ or k_2
 - then the search key x is in the tree
- Search by simply moving down the tree
- Time complexity = ?




$O(\log n)$

2-3 Tree: Insert

- Find a node L
 - One of whose children would contain the key if it were in the tree
- If L is not full
 - Then the new key is added
- If L is already full
 - **Split** L into two nodes L and L' (dividing the keys evenly among the two nodes)
 - **Promote** a copy of the *least-valued key* in L'
- Promotion may cause the parent to split in turn
 - Perhaps eventually leading to splitting the root (causing the 2-3 tree to gain a new level)


2-3 Tree: Insert - Example

- Insert a record with key value 18
 - Find the node L , one of whose children would contain 18 if it were in the tree
 - Since the parent L is not full, the key 18 can be added with no further modification to the tree




2-3 Tree: Insert - Example

- Insert a record with key value 18
 - Find the node L , one of whose children would contain 18 if it were in the tree
 - Since the parent L is not full, the key 18 can be added with no further modification to the tree
 - Update the key values in the parents




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Find the node L , one of whose children would contain 10 if it were in the tree
 - Since the parent L is already full, it has to be **split** (into L & L')




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Find the node L , one of whose children would contain 10 if it were in the tree
 - Since the parent L is already full, it has to be **split** (into L & L')
 - Next **rearrange** the keys evenly among the two nodes
 - **Update** the key values in the parents




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Find the node L , one of whose children would contain 10 if it were in the tree
 - Since the parent L is already full, it has to be split (into L & L')
 - Next rearrange the keys evenly among the two nodes
 - Update the key values in the parents
 - Promote a copy of the *least-valued key 10* in L'




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - **Promote** a copy of the *least-valued key 10* in L'
 - Grandparent is already full
 - So the *promoted key 10* cannot be inserted
 - **Split-and-promote** process *again*
 - Split L into L & L'




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Split-and-promote process again
 - Since L is the root, a **new root** has to be created as well




2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Split-and-promote process again
 - Next rearrange the keys evenly among the two nodes
 - Update the key values in the parents



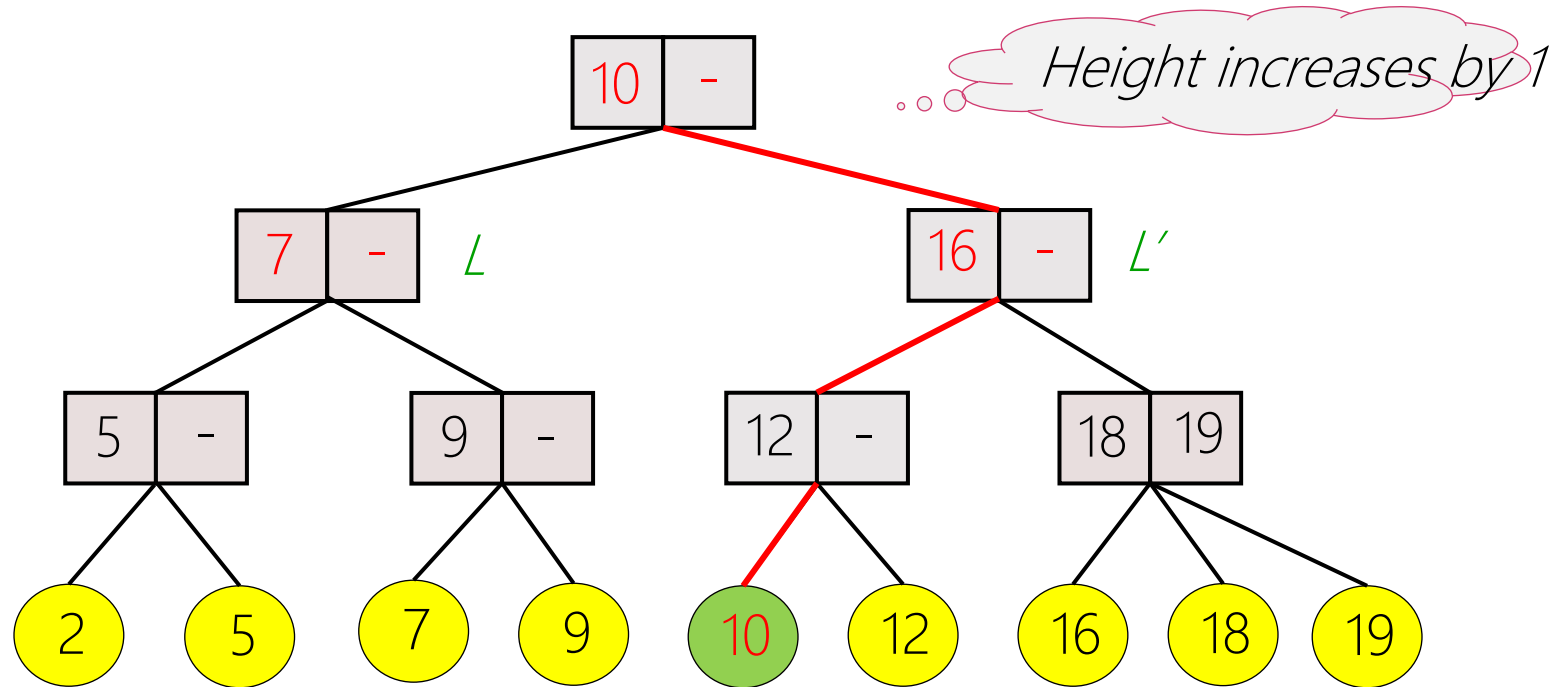
2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Split-and-promote process again
 - Promote a copy of the least-valued key 10 in L'



2-3 Tree: Insert - Example

- Insert a record with key value 10
 - Split-and-promote process again
 - The insertion is complete




2-3 Tree: Delete

- Locate *node* N , one of whose children is *key* K to be deleted
- Three cases to consider:
 - If N has 3 children
 - Just remove K
 - If N has 2 children (Underflow)
 - Look at N 's adjacent siblings to determine if they have a spare key that can be used to fill the gap
 - If not (i.e., neither sibling can lend a key to the underflow node N)
 - N must give its key to a sibling and be removed from the tree


2-3 Tree: Delete (Case-1: Three Children)

- Delete a record with key value 9
 - Find the *node N*, one of whose children is *key 9* to be deleted
 - Remove the record




2-3 Tree: Delete (Case-1: Three Children)

- Delete a record with key value 9
 - Update the key value of the parent node N




2-3 Tree: Delete (Case-1: Three Children)

- Delete a record with key value 9
 - Update the key value of the parent node N
 - The deletion is complete




2-3 Tree: Delete (Case-2: Two Children)

- Delete a record with key value 5
 - Find the *node N*, one of whose children is *key 5* to be deleted
 - Remove the record whose key is 5




2-3 Tree: Delete (Case-2: Two Children)

- Delete a record with key value 5
 - Deleting a record causes the parent node N to **underflow**
 - Look for N 's **adjacent sibling** that can lend a spare key to it




2-3 Tree: Delete (Case-2: Two Children)

- Delete a record with key value 5
 - Transfer the spare record to the underflow node N




2-3 Tree: Delete (Case-2: Two Children)

- Delete a record with key value 5
 - Transfer the spare record to the underflow node N
 - Update the key value of the parent node
 - The deletion is complete




2-3 Tree: Delete (Case-3: No Siblings with a Spare Record)

- Delete a record with key value 7
 - Find the *node N*, one of whose children is *key 7* to be deleted
 - Remove the record whose key is 7




2-3 Tree: Delete (Case-3: No Siblings with a Spare Record)

- Delete a record with key value 7
 - Deleting a record causes the parent node N to **underflow**
 - No adjacent siblings that have a spare record
 - Transfer the record (*key 9*) to a sibling




2-3 Tree: Delete (Case-3: No Siblings with a Spare Record)

- Delete a record with key value 7
 - Remove the parent node N recursively




2-3 Tree: Delete (Case-3: No Siblings with a Spare Record)

- Delete a record with key value 7
 - Remove the parent node N recursively
 - Node-merge deletion process



2-3 Tree: Delete (Case-3: No Siblings with a Spare Record)

- Delete a record with key value 7
 - Remove the parent node N recursively
 - Node-merge deletion process




B-tree (of order m)

- m -way search tree satisfying:

Shape Property

- Root has at least 2 children
- All internal nodes except root have at least $\lceil m/2 \rceil$ children
- All leaf nodes are at the same level

B-tree: Example

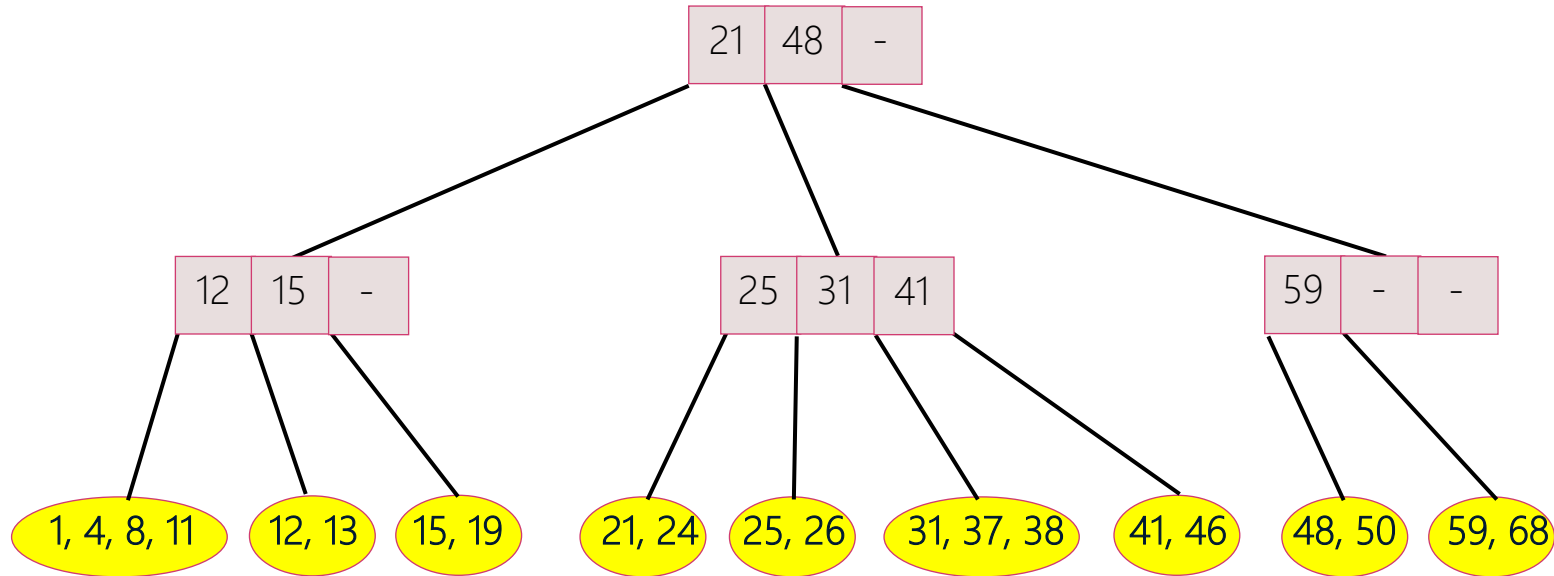


- B-tree of order 4
 - 2-3-4 tree
 - Search tree property + Shape property
- Height Bound = $O(\log n)$

B-tree

- Analysis
 - Insert, Delete, Search: $O(\log n)$
 - where the log base is the average branching factor of the tree
- Implementation of B-tree node
 - Normally equivalent to a disk block
 - Typically allows 100 or more children
 - B-tree and its variants are extremely shallow

B⁺-tree (of order 4)



- Actual records are stored only at the leaf nodes
 - A leaf node may store more or less than m (generally, between $\lceil m/2 \rceil$ & m) actual records

B*-tree

- B⁺-tree variants
 - Each node is at least 2/3 full

Comparison of B-trees with Sorted Arrays

- Disadvantages
 - More space overhead
 - More complex implementation
- Advantages
 - Updates are much easier
 - Search time is generally faster



References

- Further reading list and references
 - <https://en.wikipedia.org/wiki/B-tree>
 - <https://en.wikipedia.org/wiki/2%E2%80%93tree>
- Slide credit
 - Jaesik Park
 - Seung-Hwan Baek
 - Jong-Hyeok Lee